

homework_4_kanutala_santosh

Santosh Kanutala

7/1/2018

****My Github repository for my assignments can be found at below URL: (<https://github.com/santumagic/compscix-415-2assignments.git>)****

```
library(tidyverse)
library(mdsr)
library(nycflights13)
```

Section 5.6.7: #2, #4 and #6 only. Extra Credit: Do #5

QUESTION 2:

```
# First lets find the not cancelled flights.
not_cancelled <- flights %>%
  filter(!is.na(dep_delay), !is.na(arr_delay))

# given code for not_cancelled %>% count(dest)
not_cancelled %>%
  count(dest)
```

```
## # A tibble: 104 x 2
##   dest      n
##   <chr> <int>
## 1 ABQ    254
## 2 ACK    264
## 3 ALB    418
## 4 ANC      8
## 5 ATL  16837
## 6 AUS   2411
## 7 AVL    261
## 8 BDL    412
## 9 BGR    358
## 10 BHM   269
## # ... with 94 more rows
```

```
# new code for not_cancelled %>% count(dest) by group by and summarise
not_cancelled %>%
  group_by(dest) %>%
  summarise(n = n())
```

```
## # A tibble: 104 x 2
##   dest      n
##   <chr> <int>
## 1 ABQ    254
## 2 ACK    264
## 3 ALB    418
## 4 ANC      8
```

```
## 5 ATL 16837
## 6 AUS 2411
## 7 AVL 261
## 8 BDL 412
## 9 BGR 358
## 10 BHM 269
## # ... with 94 more rows
```

```
# given code for not_cancelled %>% count(tailnum, wt = distance)
not_cancelled %>%
  count(tailnum, wt = distance)
```

```
## # A tibble: 4,037 x 2
##   tailnum      n
##   <chr>    <dbl>
## 1 D942DN    3418
## 2 NOEGMQ 239143
## 3 N10156 109664
## 4 N102UW 25722
## 5 N103US 24619
## 6 N104UW 24616
## 7 N10575 139903
## 8 N105UW 23618
## 9 N107US 21677
## 10 N108UW 32070
## # ... with 4,027 more rows
```

```
# new code for not_cancelled %>% count(tailnum, wt = distance) group by and summarise
not_cancelled %>%
  group_by(tailnum) %>%
  summarize(n = sum(distance, na.rm = TRUE))
```

```
## # A tibble: 4,037 x 2
##   tailnum      n
##   <chr>    <dbl>
## 1 D942DN    3418
## 2 NOEGMQ 239143
## 3 N10156 109664
## 4 N102UW 25722
## 5 N103US 24619
## 6 N104UW 24616
## 7 N10575 139903
## 8 N105UW 23618
## 9 N107US 21677
## 10 N108UW 32070
## # ... with 4,027 more rows
```

QUESTION 4:

```
# number of cancelled flights per day
(cancelled_flights <- flights %>%
  filter(is.na(dep_time)) %>%
  count(day))
```

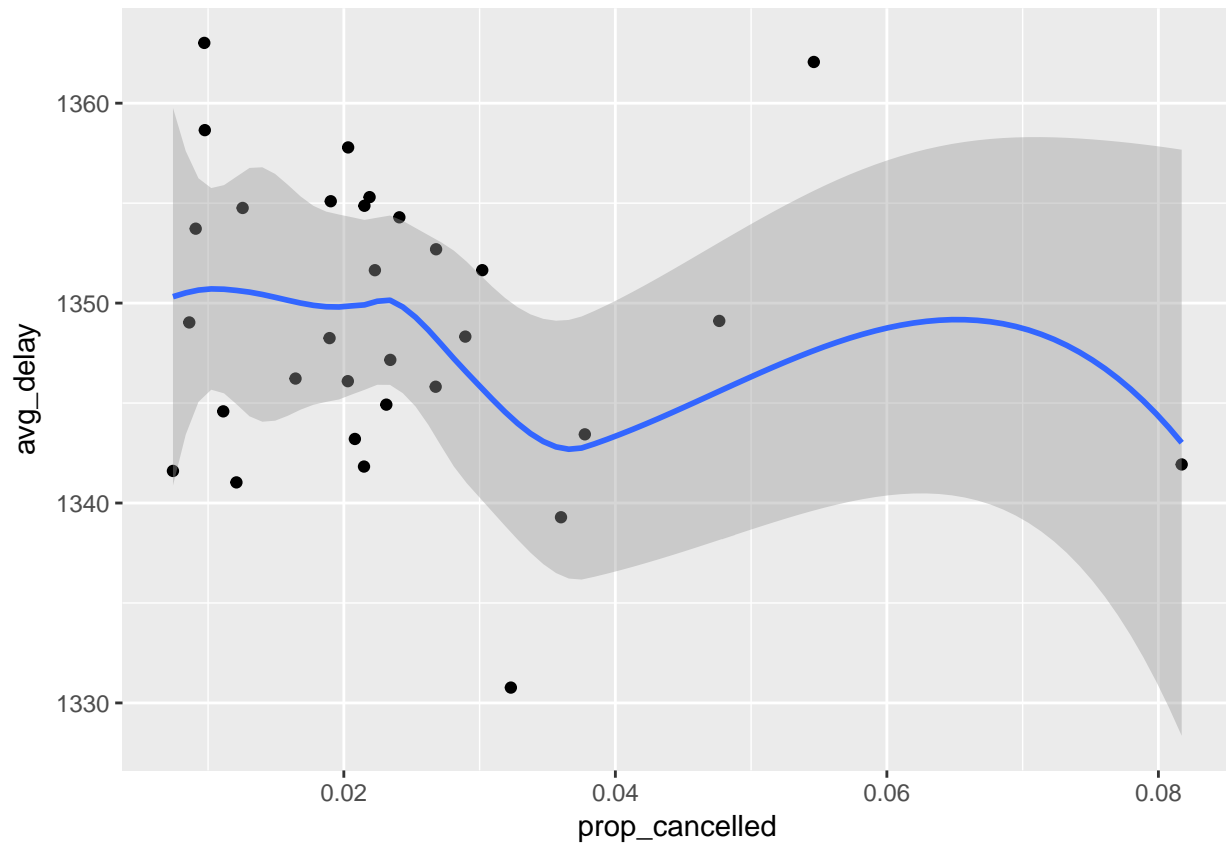
```
## # A tibble: 31 x 2
```

```
##      day      n
##    <int> <int>
##  1      1    246
##  2      2    250
##  3      3    109
##  4      4     82
##  5      5    226
##  6      6    296
##  7      7    318
##  8      8    921
##  9      9    593
## 10     10    535
## # ... with 21 more rows

# proportion of day cancelled flights vs aerage delays
(cancelled_flights <- flights %>%
  group_by(day) %>%
  summarize(prop_cancelled = sum(is.na(dep_time)) / n(),
            avg_delay = mean(dep_time, na.rm = TRUE)))
```

```
## # A tibble: 31 x 3
##       day prop_cancelled avg_delay
##   <int>         <dbl>     <dbl>
## 1     1           0.0223    1352.
## 2     2           0.0231    1345.
## 3     3           0.00972   1363.
## 4     4           0.00741   1342.
## 5     5           0.0208    1343.
## 6     6           0.0268    1346.
## 7     7           0.0289    1348.
## 8     8           0.0817    1342.
## 9     9           0.0546    1362.
## 10    10          0.0477    1349.
## # ... with 21 more rows
```

```
# plot for the relationship
ggplot(cancelled_flights, aes(x = prop_cancelled, y = avg_delay)) +
  geom_point() +
  geom_smooth()
```



QUESTION 6:

ANSWER 6:

sort argument will sort the elements of count () in the descending order. After the results are extracted, we can use the sort to arrange the values.

Section 10.5: #1, #2, #3 and #6 only.

QUESTION 2:

```
# mtcars as a data frame
print(mtcars)
```

```
##          mpg  cyl  disp  hp drat   wt  qsec vs  am  gear  carb
## Mazda RX4      21.0   6  160.0  110 3.90 2.620 16.46 0   1    4    4
## Mazda RX4 Wag  21.0   6  160.0  110 3.90 2.875 17.02 0   1    4    4
## Datsun 710      22.8   4  108.0   93 3.85 2.320 18.61 1   1    4    1
## Hornet 4 Drive  21.4   6  258.0  110 3.08 3.215 19.44 1   0    3    1
## Hornet Sportabout 18.7   8  360.0  175 3.15 3.440 17.02 0   0    3    2
## Valiant         18.1   6  225.0  105 2.76 3.460 20.22 1   0    3    1
## Duster 360      14.3   8  360.0  245 3.21 3.570 15.84 0   0    3    4
## Merc 240D       24.4   4  146.7   62 3.69 3.190 20.00 1   0    4    2
## Merc 230        22.8   4  140.8   95 3.92 3.150 22.90 1   0    4    2
## Merc 280        19.2   6  167.6  123 3.92 3.440 18.30 1   0    4    4
```

```
## Merc 280C      17.8   6 167.6 123 3.92 3.440 18.90 1 0   4   4
## Merc 450SE     16.4   8 275.8 180 3.07 4.070 17.40 0 0   3   3
## Merc 450SL     17.3   8 275.8 180 3.07 3.730 17.60 0 0   3   3
## Merc 450SLC    15.2   8 275.8 180 3.07 3.780 18.00 0 0   3   3
## Cadillac Fleetwood 10.4  8 472.0 205 2.93 5.250 17.98 0 0   3   4
## Lincoln Continental 10.4  8 460.0 215 3.00 5.424 17.82 0 0   3   4
## Chrysler Imperial 14.7  8 440.0 230 3.23 5.345 17.42 0 0   3   4
## Fiat 128       32.4   4  78.7  66 4.08 2.200 19.47 1 1   4   1
## Honda Civic    30.4   4  75.7  52 4.93 1.615 18.52 1 1   4   2
## Toyota Corolla 33.9   4  71.1  65 4.22 1.835 19.90 1 1   4   1
## Toyota Corona  21.5   4 120.1  97 3.70 2.465 20.01 1 0   3   1
## Dodge Challenger 15.5  8 318.0 150 2.76 3.520 16.87 0 0   3   2
## AMC Javelin    15.2  8 304.0 150 3.15 3.435 17.30 0 0   3   2
## Camaro Z28     13.3  8 350.0 245 3.73 3.840 15.41 0 0   3   4
## Pontiac Firebird 19.2  8 400.0 175 3.08 3.845 17.05 0 0   3   2
## Fiat X1-9      27.3   4  79.0  66 4.08 1.935 18.90 1 1   4   1
## Porsche 914-2  26.0   4 120.3  91 4.43 2.140 16.70 0 1   5   2
## Lotus Europa   30.4   4  95.1 113 3.77 1.513 16.90 1 1   5   2
## Ford Pantera L 15.8  8 351.0 264 4.22 3.170 14.50 0 1   5   4
## Ferrari Dino   19.7   6 145.0 175 3.62 2.770 15.50 0 1   5   6
## Maserati Bora   15.0  8 301.0 335 3.54 3.570 14.60 0 1   5   8
## Volvo 142E     21.4   4 121.0 109 4.11 2.780 18.60 1 1   4   2
```

```
# mtcars as a tibble
print(as_tibble(mtcars))
```

```
## # A tibble: 32 x 11
##   mpg   cyl  disp    hp  drat    wt  qsec    vs  am  gear  carb
##   <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
## 1  21     6   160   110   3.9   2.62  16.5     0     1     4     4
## 2  21     6   160   110   3.9   2.88  17.0     0     1     4     4
## 3 22.8     4   108    93   3.85   2.32  18.6     1     1     4     1
## 4 21.4     6   258   110   3.08   3.22  19.4     1     0     3     1
## 5 18.7     8   360   175   3.15   3.44  17.0     0     0     3     2
## 6 18.1     6   225   105   2.76   3.46  20.2     1     0     3     1
## 7 14.3     8   360   245   3.21   3.57  15.8     0     0     3     4
## 8 24.4     4   147.    62   3.69   3.19  20.0     1     0     4     2
## 9 22.8     4   141.    95   3.92   3.15  22.9     1     0     4     2
## 10 19.2     6   168.   123   3.92   3.44  18.3     1     0     4     4
## # ... with 22 more rows
```

When we print a dataframe whole records will be printed where as with the tibble only top rows will be printed and the columns and widths are adjusted to fit the screen resolution.

QUESTION 2:

```
# with a data frame
df <- data_frame(abc = 1, xyz = 'a')
df

## # A tibble: 1 x 2
##   abc xyz
##   <dbl> <chr>
## 1     1 a
```

```
df$x

## NULL

df[, "xyz"]

## # A tibble: 1 x 1
##   xyz
##   <chr>
## 1 a

df[, c("abc", "xyz")]

## # A tibble: 1 x 2
##   abc xyz
##   <dbl> <chr>
## 1     1 a

# with a tibble
df <- tibble(abc = 1, xyz = "a")
df

## # A tibble: 1 x 2
##   abc xyz
##   <dbl> <chr>
## 1     1 a

df$x

## NULL

df[, "xyz"]

## # A tibble: 1 x 1
##   xyz
##   <chr>
## 1 a

df[, c("abc", "xyz")]

## # A tibble: 1 x 2
##   abc xyz
##   <dbl> <chr>
## 1     1 a
```

By using \$ we can only extract by names. tibbles are much stricter and they alert us when we are trying to pull a column which never existed.

QUESTION 3:

```
var <- "mpg"
```

ANSWER:

We can extract the reference variable by using double square brackets like `df[[var]]`

QUESTION 6:

```
?tibble

## Help on topic 'tibble' was found in the following packages:
##
##   Package          Library
##   tibble           /Library/Frameworks/R.framework/Versions/3.5/Resources/library
##   dplyr            /Library/Frameworks/R.framework/Versions/3.5/Resources/library
##
##
## Using the first match ...
```

ANSWER:

```
getOption("tibble.max_extra_cols")

## NULL
```

Section 12.3.3: #2, #3 and #4 only.

QUESTION 2:

```
# table4a %>%
# gather(1999, 2000, key = "year", value = "cases") # commented out as it is failing and need to ask
```

ANSWER:

we need to add the quotes around 1999 & 2000 otherwise 1999 & 2000 are treated as columns in the data frame. Below is the modified code.

```
table4a %>%
  gather('1999', '2000', key = "year", value = "cases")

## # A tibble: 6 x 3
##   country    year  cases
##   <chr>      <chr> <int>
## 1 Afghanistan 1999     745
## 2 Brazil      1999   37737
## 3 China       1999  212258
## 4 Afghanistan 2000    2666
## 5 Brazil      2000   80488
## 6 China       2000  213766
```

QUESTION 3:

```
# given code
people <- tribble(
  ~name,      ~key,      ~value,
  #-----/-----/-----
```

```

"Phillip Woods", "age", 45,
"Phillip Woods", "height", 186,
"Phillip Woods", "age", 50,
"Jessica Cordero", "age", 37,
"Jessica Cordero", "height", 156
)

# Spreading the above tibble
# spread(people, key, value) # commented out as it is failing and need to ask instructor

```

ANSWER:

Duplicates are identified. In other words Phillip has two age records. In this case we can add another column related to the row identifier like the `people_id` then the issue will be resolved. Below is the code for the same.

```

people$people_id <- c(1, 1, 2, 1, 1)
people # display the modified tibble

## # A tibble: 5 x 4
##   name          key    value people_id
##   <chr>         <chr> <dbl>     <dbl>
## 1 Phillip Woods age      45         1
## 2 Phillip Woods height  186         1
## 3 Phillip Woods age      50         2
## 4 Jessica Cordero age      37         1
## 5 Jessica Cordero height  156         1

spread(people, key, value) # spread the tibble

```

```

## # A tibble: 3 x 4
##   name          people_id age height
##   <chr>         <dbl> <dbl> <dbl>
## 1 Jessica Cordero      1    37    156
## 2 Phillip Woods      1    45    186
## 3 Phillip Woods      2    50     NA

```

QUESTION 4:

```

preg <- tribble(
  ~pregnant, ~male, ~female,
  "yes",     NA,    10,
  "no",      20,    12
)

```

ANSWER:

By looking at above tibble, we can observe that there is a missing value. So we can use `gather()` and add `na.rm = TRUE` to remove the value NA. Below is the gathered tibble with variables, pregnant, gender, count.

```

preg %>%
  gather(key = 'gender', value = 'value', c(2:3), na.rm = TRUE)

```



```
## # A tibble: 3 x 3
##   pregnant gender value
## * <chr>    <chr> <dbl>
## 1 no      male    20
## 2 yes     female   10
## 3 no      female   12
```

Section 12.4.3: #1 and #2 only.

QUESTION 1:

```
tibble(x = c("a,b,c", "d,e,f,g", "h,i,j")) %>%
  separate(x, c("one", "two", "three"))
```

```
## # A tibble: 3 x 3
##   one  two  three
##   <chr> <chr> <chr>
## 1 a    b    c
## 2 d    e    f
## 3 h    i    j
```

```
tibble(x = c("a,b,c", "d,e", "f,g,i")) %>%
  separate(x, c("one", "two", "three"))
```

```
## # A tibble: 3 x 3
##   one  two  three
##   <chr> <chr> <chr>
## 1 a    b    c
## 2 d    e   <NA>
## 3 f    g    i
```

```
# help on the separate
?separate
```

ANSWER:

‘extra’ argument controls what happens when there are too many pieces. It can drop extra values by warning or without warning or it can merge.

‘fill’ controls what happens when there are not enough piece by filling the missing values from right or left with or without a warning.

```
# using extra = "drop"
tibble(x = c("a,b,c", "d,e,f,g", "h,i,j")) %>%
  separate(x, c("one", "two", "three"), extra = "drop")
```

```
## # A tibble: 3 x 3
##   one  two  three
##   <chr> <chr> <chr>
## 1 a    b    c
## 2 d    e    f
## 3 h    i    j
```

```
# using extra = "merge"
tibble(x = c("a,b,c", "d,e,f,g", "h,i,j")) %>%
  separate(x, c("one", "two", "three"), extra = "merge")
```

```
## # A tibble: 3 x 3
##   one  two  three
##   <chr> <chr> <chr>
## 1 a    b    c
## 2 d    e    f,g
## 3 h    i    j
```

```
# using fill = "right"
tibble(x = c("a,b,c", "d,e", "f,g,i")) %>%
  separate(x, c("one", "two", "three"), fill = "right")
```

```
## # A tibble: 3 x 3
##   one  two  three
##   <chr> <chr> <chr>
## 1 a    b    c
## 2 d    e    <NA>
## 3 f    g    i
```

```
# using fill = "left"
tibble(x = c("a,b,c", "d,e", "f,g,i")) %>%
  separate(x, c("one", "two", "three"), fill = "left")
```

```
## # A tibble: 3 x 3
##   one  two  three
##   <chr> <chr> <chr>
## 1 a    b    c
## 2 <NA> d    e
## 3 f    g    i
```

QUESTION 2:

Both unite() and separate() have a remove argument. What does it do? Why would you set it to FALSE?

ANSWER:

remove argument will remove the original input columns from the output. If we change it to FALSE then the original columns are retained in the output.

Data import Exercises:

```
# Load the baby_names file
baby_names_txt <- read.delim("/Users/skanutal/Documents/Santosh/Learning/Berkeley/baby_names.txt", sep=
# glimpse the data from txt file
glimpse(baby_names_txt)
```

```
## Observations: 30,000
## Variables: 5
## $ year <int> 1880, 1880, 1880, 1880, 1880, 1880, 1880, 1880, 1880, 188...
## $ sex <fct> F, F, F, F, F, F, F, F, F, F, F, F, F, F, F, F, F, F, F, ...
```

```
## $ name <fct> Mary, Anna, Emma, Elizabeth, Minnie, Margaret, Ida, Alice...
## $ n    <int> 7065, 2604, 2003, 1939, 1746, 1578, 1472, 1414, 1320, 128...
## $ prop <dbl> 0.072384329, 0.026679234, 0.020521700, 0.019865989, 0.017...

# Export the .txt file as csv into a different location
baby_names_csv <- write_delim(
  baby_names_txt, delim = '|', path = "/Users/skanutal/Documents/Santosh/Learning/Berkeley/3. Intro to D
)
# Load the newly created csv file
baby_names_csv_final <- read.csv(
  "/Users/skanutal/Documents/Santosh/Learning/Berkeley/3. Intro to DS/Assignments/baby_names.csv", sep=
)
# glimpse the data from the final csv file
glimpse(baby_names_csv_final)

## Observations: 30,000
## Variables: 5
## $ year <int> 1880, 1880, 1880, 1880, 1880, 1880, 1880, 1880, 1880, 188...
## $ sex  <fct> F, F, F, F, F, F, F, F, F, F, F, F, F, F, F, F, F, F, ...
## $ name <fct> Mary, Anna, Emma, Elizabeth, Minnie, Margaret, Ida, Alice...
## $ n    <int> 7065, 2604, 2003, 1939, 1746, 1578, 1472, 1414, 1320, 128...
## $ prop <dbl> 0.072384329, 0.026679234, 0.020521700, 0.019865989, 0.017...
```